

Getting Data into R

Brendan Apfeld

October 14, 2016

A Quick R Refresher

Packages

- R is super useful, but we supplement its functionality with *packages*
- You install packages (assuming they're in the normal place) using `install.packages("package_name")`
 - You only have to install a package once (though you may need to update it later)
- You have to *load* packages every time you start an R session
 - Load a package in R by using `library(package_name)`
 - Note there are no quotation marks there
 - Loading a package makes all of its functions available to you
 - If you don't need that, you can also do `package_name::function()`
 - E.g. (that we'll see in a second) `rio::import("file_name.csv")`

Objects and Operators

- Pretty much *everything* in R is an object
 - Which is probably why it's an “OO,” or “object-oriented” language
- You “assign” objects to names in R using the `<-` operator

Data Sources

Online

- Sometimes you'll have a direct link to a csv or txt file
- Other online data may be accessed through an api (and an R package)
- Advantages: potentially easy, “cutting edge data”
- Disadvantages: might change/disappear, often slow

On Disk

- More often, you'll have some kind of data file on disk
- Advantages: fast, it shouldn't disappear
- Disadvantages: you can break it

Data Types

Common Data File Formats

- Spreadsheet-style
 - .xlsx or .xls (Excel)
 - .ods (OpenDocument)
 - Google Sheets
- Proprietary Formats
 - .dta (Stata)
 - .sav (SPSS)
 - .xpt (SAS)
 - .Rdata (R)
- Plain Text Data Formats
 - .csv (our good friend)
 - .dat (not very common)
 - .json
 - .xml

Getting Data into R

The (Bad) Old Days

- `read.csv()` or `readr::read_csv()`
- `gdata::read.xls()`
- `foreign::read.dta()` or `haven::read_dta()`
- Which one should I use?
 - Depended on the version of the program used to create the data
 - Each package required different arguments
 - You never really knew what to expect

One Import Package to Rule them All

Then the `rio` package came along

- Only two commands in the whole package: `rio::import()` and `rio::export()`
- The only argument you *need* for it to work is the file name
- Supports other arguments if you really *want* to change its defaults
- You (almost) never need to change the defaults

Saving and Opening Rdata Files

- R has its own data format that ends in .Rdata
 - Super compressed
 - Saves and loads quickly
- Save with `save(objects, file = "filename.Rdata")`
- Load with `load("filename.Rdata")`

Examples with Data on Disk

Examples in R

Examples with Data Online

- Can be as simple as

```
wunderground <-  
read.csv("https://www.wunderground.com/history/airport/ZBAA/  
2013/1/1/DailyHistory.html?format=1")
```

- The list of possibilities for APIs is far to great to include here
- For an example using the World Bank data API and the corresponding `wbstats` package, check out my shameless self-promotion

Visualizing Data

Visualizing Datasets

- Datasets are often large and unwieldy
- There are a number of ways to easily learn something about them

```
mtcars <- rio::import("mtcars.csv")
mtcars # a lot of times this is too much for the console
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## 1	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
## 2	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
## 3	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
## 4	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## 5	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
## 6	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
## 7	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
## 8	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
## 9	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2

Visualizing Datasets

```
str(mtcars) # look at the structure of the data
```

```
## 'data.frame':    32 obs. of  11 variables:
##  $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
##  $ cyl : int   6 6 4 6 8 6 8 4 4 6 ...
##  $ disp: num  160 160 108 258 360 ...
##  $ hp  : int  110 110 93 110 175 105 245 62 95 123 ...
##  $ drat: num   3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
##  $ wt  : num   2.62 2.88 2.32 3.21 3.44 ...
##  $ qsec: num   16.5 17 18.6 19.4 17 ...
##  $ vs  : int    0 0 1 1 0 1 0 1 1 1 ...
##  $ am  : int    1 1 1 0 0 0 0 0 0 0 ...
##  $ gear: int    4 4 4 3 3 3 3 4 4 4 ...
##  $ carb: int    4 4 1 1 2 1 4 2 2 4 ...
```

Visualizing Datasets

```
names(mtcars) # get the names of the variables in the dataframe
```

```
## [1] "mpg"  "cyl"  "disp" "hp"   "drat" "wt"   "qsec" "vs"   "am"   "
## [11] "carb"
```

```
head(mtcars) # print the first 6 lines
```

```
##      mpg  cyl disp  hp drat    wt  qsec vs  am gear carb
## 1  21.0    6  160  110  3.90  2.620  16.46  0   1    4    4
## 2  21.0    6  160  110  3.90  2.875  17.02  0   1    4    4
## 3  22.8    4  108   93  3.85  2.320  18.61  1   1    4    1
## 4  21.4    6  258  110  3.08  3.215  19.44  1   0    3    1
## 5  18.7    8  360  175  3.15  3.440  17.02  0   0    3    2
## 6  18.1    6  225  105  2.76  3.460  20.22  1   0    3    1
```

R Also has a Viewer

```
wg <- read.csv("https://www.wunderground.com/history/airport/ZBAA/2013/1")  
# View(wg) # commented out because it won't run on a slide  
# note the capital "V"  
# Can also access by clicking on the object in the  
# environment pane in RStudio
```

Graphing - Back to R!